

Master in Artificial Intelligence

Document
structure

Language
identification

Introduction to Human Language Technologies



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Facultat d'Informàtica de Barcelona



Outline

Document
structure

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Outline

Document
structure

Searching
textual zones

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting
- 2 Language identification

Document types

Document
structure

Searching
textual zones

Language
identification

- Documents containing text:
 - Structured documents (e.g., web pages being tables)
 - Semi-structured documents (e.g., web pages containing pieces of plain text, figures and tables)
 - Documents with plain text only (e.g., text files, emails, tweets, oral transcripts)

Accessing to plain text contained in web pages may be relevant.

XML Parsers

- Transform an XML/HTML/XHTML document into a tree of standard objects.
- Provide an interface to manage that tree.
- Textual zones in the document can be extracted from that tree using the interface.

Document
structure

Searching
textual zones

Language
identification

```
<?xml version="1.0"?>
<doc type="novel" title="The green apple">
<chapter id="1">
<p>There are lots of trees in Amsteel Hill. I remember
going there and spend all the morning climbing those
trees, trying to get as many apples as possible.</p>
<p> James always wanted to come with me but he
was too young to get climbing.</p>
...
</doc>
```

Using ElementTree.py

```
for c in root:
    lp=c.findall('p')
    for p in lp:
        print p.text
```

Outline

Document
structure
Tokenization
Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting
- 2 Language identification

Goal of tokenization

- Goal: split plain text into *basic units*
- Use: IR tasks, text categorization, sentence splitting, language identification, text normalization . . .
- Different *basic units* depending on the task,
 - *Naïve* tokenizations: split by blanks and punctuation marks occurring after alphanum-string.
 - Complex tokenizations: names, clitics, abbreviations, *collocations* . . .

Goal of tokenization

- Goal: split plain text into *basic units*
- Use: IR tasks, text categorization, sentence splitting, language identification, text normalization . . .
- Different *basic units* depending on the task,
 - Naïve tokenizations: split by blanks and punctuation marks occurring after alphanum-string.
 - Complex tokenizations: names, clitics, abbreviations, *collocations* . . .

Relevant definitions:

Word N-gram: sequence of words occurring in a text

Collocation: sequence of words that frequently occur together. Ex: "break a leg", "On the one hand"

Examples of tokenization

Blanks	outer punct.	Abbr.	Clitics	Colloc.	text normalized
Of	Of	Of	Of	Of_course	Of_course
course	course	course	course		
I'll	I'll	I'll	I	I	I
			'll	'll	will
go	go	go	go	go	go
to	to	to	to	to	to
U.P.C.	U.P.C	U.P.C	U.P.C	U.P.C	Universitat. . .
	:	:	:	:	:
"Daily,	Daily	Daily	Daily	Daily	Daily
Mr.	Mr	Mr.	Mr.	Mr.	Mister
John	John	John	John	John	John_Smith
Smith..."	Smith	Smith	Smith	Smith	

	"	"	"	"	"

Document
structure
Tokenization
Language
identification

Examples of tokenization

Blanks	outer punct.	Abbr.	Clitics	Colloc.	text normalized
Of	Of	Of	Of	Of_course	Of_course
course	course	course	course		
I'll	I'll	I'll	I	I	I
			'll	'll	will
go	go	go	go	go	go
to	to	to	to	to	to
U.P.C.	U.P.C	U.P.C	U.P.C	U.P.C	Universitat. . .
	„	„	„	„	„
"Daily,	Daily	Daily	Daily	Daily	Daily
	,	,	,	,	,
Mr.	Mr	Mr.	Mr.	Mr.	Mister
	.				
John	John	John	John	John	John_Smith
Smith..."	Smith	Smith	Smith	Smith	

	„	„	„	„	„

Document
structure
Tokenization
Language
identification

Examples of tokenization

Blanks	outer punct.	Abbr.	Clitics	Colloc.	text normalized
Of	Of	Of	Of	Of_course	Of_course
course	course	course	course		
I'll	I'll	I'll	I	I	I
			'll	'll	will
go	go	go	go	go	go
to	to	to	to	to	to
U.P.C.	U.P.C	U.P.C	U.P.C	U.P.C	Universitat. . .
	,	,	,	,	,
"Daily,	Daily	Daily	Daily	Daily	Daily
	,	,	,	,	,
Mr.	Mr	Mr.	Mr.	Mr.	Mister
	.				
John	John	John	John	John	John_Smith
Smith..."	Smith	Smith	Smith	Smith	

	,	,	,	,	,

Problems: Non-standard text? Chinese? Japanese?

Outline

Document
structure

Sentence
splitting

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Goal of sentence splitting

- Goal: Recognition of sentence boundaries in plain text (e.g., '. ' '? ' '! '...').
- Language-dependent task
 - Ex: German: "Mein 2. Semester kommt bald zu Ende."
 - Ex: Traditional chinese?
- Domain-dependent task
 - Ex: "It is expressed as $(x=1)$? T.add('-') : T.add(x)."
- Methods:
 - Hand-crafted rules
 - Machine learning methods
- Input:
 - Naïve tokenization that depends on the particular method.
 - For simplicity, we will assume *blanks+outer_punctuation*
 - " I'll go to U.P.C. " Daily, Mr. John Smith..." "
 - " I 'll go to U.P.C . " Daily , Mr . John Smith ... " "

Problems of sentence splitting

Main problems:

- Abbreviations and acronyms (most difficult one)

Ex: "I will meet with Mr. Smith to talk about it."

Ex: "Lisa run 25 km. She ended up in N.Y."

How to detect them?

- Ellipsis

Ex: "There're different methods (A, B, ...) but ..."

- Internal quotation

Ex: " 'Stop!' he shouted."

- Ordinal numbers (German)

- Special cases:

Ex: " We have some variables. x stands for the weight,"

Hand-crafted rules for sentence splitting

- Specific hand-crafted rules for specific cases
 - Abbreviation classes (Lists of abbreviations)
(month name, unit-of-measure, title, address name, ...)
Ex: TITLE=('Mr', 'Mrs', 'Dr', ...)
 - Regular expressions for general cases, abbreviations, ellipsis, ...
Ex: / ([?!])+ / $\rightarrow t \in \text{s_boundary}$
Ex: / (\.)\{3\} [A-Z] / $\rightarrow t \in \text{s_boundary}$
Ex: / [?!.] \) [A-Z] / $\rightarrow t \in \text{s_boundary}$
Ex: / (\$TITLE) \. / $\rightarrow t \notin \text{s_boundary}$
Ex: / [A-Z] \. / $\rightarrow t \notin \text{s_boundary}$
- Problem:
 - Highly expensive adaptation to new languages
(rules and abbreviation classes)

Supervised ML for sentence splitting

- The most frequently used (ME, SVM, CRF, ...)
- Require manually annotated corpora. Commonly,
 $e^+, e^- = [',', '!', '?']$ and some preceding and following tokens
- Represent each e as a set of features. Depends on the approach, the language and the domain, although normally they tend to be binary features.
- Problem:
 - Require very large sets of examples (tens of thousands to hundreds of thousands)

Supervised ML for sentence splitting

- Examples of features used in the state of the art
 - tok-1_X: 1st token before '.' is X
 - tok-2_X: 2nd token before '.' is X
 - tok+1_X: 1st token after '.' is X
 - len_tok-1_X: length of 1st token before '.' is X
 - len_tok-2_X: length of 2nd token before '.' is X
 - len_tok+1_X: length of 1st token after '.' is X
 - [up|lo|cap|num]_tok-1: 1st token before '.' is Upper, Lower, CAP, Numbers
 - [up|lo|cap|num]_tok-2: same for 2nd token before '.'
 - [up|lo|cap|num]_tok+1: same for 1st token after '.'
 - class_tok-1_X: abbreviation class of 1st token before '.' is X
 - ...

Supervised ML for sentence splitting

Example of annotation and binary features extraction

I 'll go to U.P.C. " Daily , Mr John Smith ... "

e^+	tok-1_U.P.C	e^-	tok-1_Mr
	len_tok-1_3		len_tok-1_2
	CAP_tok-1		up_tok-1
	tok-2_to		tok-2,
	len_tok-2_2		len_tok-2_1
	lo_tok-2		class_tok-1_title
	tok+1_"		tok+1_John
	len_tok+1_1		len_tok+1_4
			up_tok+1

Document
structure

Sentence
splitting

Language
identification

Unsupervised ML for sentence splitting

Document
structure

Sentence
splitting

Language
identification

- Based on corpus statistics
- Easily adaptable to new languages
 - They require large unannotated training corpora
- Mainly focus on abbreviations and ellipsis
- Heuristics and statistics calculated from the training corpus to decide:
 - 1 Which tokens are abbreviations?
 - 2 When the final period of the elements is a sentence boundary?
- Example: Punkt [Kiss and Strunk, 2006]

Unsupervised ML for sentence splitting

1 Punkt: Is token t considered an abbreviations?

Measured by considering the following heuristics:

- $t' = \langle t, . \rangle$ should be a collocation
- the length of t should be short
- t could include periods (acronyms)
- t is not ordinary word preceeding a period most of the times. (e.g., verbs in Turkish)

Unsupervised ML for sentence splitting

1 Punkt: Is token t considered an abbreviations?

Measured by considering the following heuristics:

- $t' = \langle t, . \rangle$ should be a collocation
- the length of t should be short
- t could include periods (acronyms)
- t is not ordinary word preceeding a period most of the times. (e.g., verbs in Turkish)

2 Punkt: Is the final period of abbreviation $t' = \langle t, . \rangle$ considered sentence boundary?

Either one of the following heuristics must be true:

- $t'' = \text{following}(t')$ is a frequent sentence (from [1]) starter
- t'' is uppercase, occurs at least once in lowercase in the training corpus but never in uppercase inside sentences (from [1])

Exercise

Document
structure

Sentence
splitting

Language
identification

Explain why Punkt fails (red) or not (blue) with the following texts:

- " "Good night!", said Laura. "
- " Abbrev. is a common abbreviation of abbreviation. "
- " We are meeting with our mr. You are late! "
- " We are meeting with our Mr. However, we'll finish soon."

Outline

Document
structure

Language
identification

- 1 Document structure
 - Searching textual zones
 - Tokenization
 - Sentence splitting

- 2 Language identification

Goal of language identification

Document
structure

Language
identification

- Can be seen as a particular classification problem.
- Given a document, d , and a set of languages, $L = \{l_1, \dots, l_k\}$, assign l_i to d .
- Method:
 - $\hat{d} = \text{representation}(d)$
 - $M(\hat{d}) \rightarrow l_i$
- Model M can be learned from training corpus $T = \{T_i\}_{1 \dots k}$ where $T_i = \{d_x | d_x \text{ written in } l_i\}$:
 - Supervised Machine Learning methods
 - Statistical Language models

Survey: <https://arxiv.org/pdf/1804.08186.pdf>

Language models for language identification

Method with language models:

$$M = \{P^{l_i}\}_{l_i \in L}$$

$P^{l_i}(\hat{d})$: probability of \hat{d} to belong to l_i

$$l_i = \operatorname{argmax}_{l \in L} (P^l(\hat{d}))$$

$P^{l_i}(\hat{d}) \approx P^{T_i}(\hat{d})$: probability of \hat{d} observing data from T_i

Language models for language identification

Method with language models:

$$M = \{P^{l_i}\}_{l_i \in L}$$

$P^{l_i}(\hat{d})$: probability of \hat{d} to belong to l_i

$$l_i = \operatorname{argmax}_{l \in L} (P^l(\hat{d}))$$

$P^{l_i}(\hat{d}) \approx P^{T_i}(\hat{d})$: probability of \hat{d} observing data from T_i

- 1 Which is the representation \hat{d} ?
- 2 How is $P^{T_i}(\hat{d})$ computed?

Language models for language identification

Method with language models:

$$M = \{P^{l_i}\}_{l_i \in L}$$

$P^{l_i}(\hat{d})$: probability of \hat{d} to belong to l_i

$$l_i = \operatorname{argmax}_{l \in L} (P^l(\hat{d}))$$

$P^{l_i}(\hat{d}) \approx P^{T_i}(\hat{d})$: probability of \hat{d} observing data from T_i

- 1 Which is the representation \hat{d} ?
- 2 How is $P^{T_i}(\hat{d})$ computed?

They depend on the particular type of model.

Most frequently used: **unigram language models**

Unigram language models for language identification

1 Which is the representation \hat{d} ?

$\hat{d} = e_1, \dots, e_s$ being the occurrences of unigrams:

- Words (after *Naïve* tokenization) or
- Characters n -grams (tokenization is not required)
 - n fixed (the most frequently used) or
 - n variable (improves accuracy, lower efficiency)

Unigram language models for language identification

1 Which is the representation \hat{d} ?

$\hat{d} = e_1, \dots, e_s$ being the occurrences of unigrams:

- Words (after *Naïve* tokenization) or
- Characters n -grams (tokenization is not required)
 - n fixed (the most frequently used) or
 - n variable (improves accuracy, lower efficiency)

2 How is $P^T(\hat{d})$ computed?

Each e_j is independent from the rest

$$P^T(\hat{d}) = P^T(e_1, \dots, e_s) = \prod_{j=1}^s P^T(e_j)$$

$$\log P^T(\hat{d}) = \sum_{j=1}^s \log P^T(e_j)$$

Possible estimators of $P^T(e_j)$:

- Maximum Likelihood Estimator (MLE)
- Smoothing techniques.

Unigram language models for language identification

Maximum Likelihood Estimator

$$P^T(e_j) \approx P_{MLE}^T(e_j) = \frac{c_T(e_j)}{N_T}$$

$c_T(x)$: #observed occurrences of x in training corpus T

N_T : #observed occurrences of elements in training corpus T

Unigram language models for language identification

Maximum Likelihood Estimator

$$P^T(e_j) \approx P_{MLE}^T(e_j) = \frac{c_T(e_j)}{N_T}$$

$c_T(x)$: #observed occurrences of x in training corpus T

N_T : #observed occurrences of elements in training corpus T

- Problem: data sparseness. Unseen e_j causes the model to fail. MLE is unsuitable for NLP.

Unigram language models for language identification

Maximum Likelihood Estimator

$$P^T(e_j) \approx P_{MLE}^T(e_j) = \frac{c_T(e_j)}{N_T}$$

$c_T(x)$: #observed occurrences of x in training corpus T

N_T : #observed occurrences of elements in training corpus T

■ Example:

$P^{[en]}('The\ doctor\ tell\ us\ about\ his\ quadriplegia')?$

$$\begin{aligned} c_{[en]}('quadriplegia') = 0 &\implies P_{MLE}^{[en]}('quadriplegia') = 0 \\ \implies P^{[en]}('The\ doctor\ tell\ us\ about\ his\ quadriplegia') &= 0 !! \end{aligned}$$

Unigram language models for language identification

Document
structure

Language
identification

Smoothing Techniques:

Keep some probability mass for e_j unseen in T_i

E.g., Lidstone's Law (LID)

$$P^T(e_j) \approx P_{LID}^T(e_j) = \frac{c_T(e_j) + \lambda}{N_T + \lambda B} \quad \text{usually, } \lambda = 0,5$$

B : #bins (potentially observable unigrams)

Exercise

Suppose we have a Language Identifier for English and Catalan, based on unigram language models with words and the following statistics

w_i	a	he	mail	sent	to	mordorian
English language model [en]						
$c_{[en]}(w_i)$	17.000	10.000	3.900	850	25.000	0
$N_{[en]}=1.300.000$	$B_{[en]}=22.600$					
Catalan Language model [ca]						
$c_{[ca]}(w_i)$	21.000	11.900	420	910	750	0
$N_{[ca]}=1.100.000$	$B_{[ca]}=36.800$					

- Compute $P^{[en]}$ and $P^{[ca]}$ using MLE and LID for the following texts:
 - "he"
 - "he sent a"
 - "he sent a mail"
 - "he sent a mail to a mordorian"
- What language is identified by each estimator for each of the previous texts?
- Explain the effects of the text size